Embedded Base Boot Requirements (EBBR) Specification

Release v2.0.0-17-g8a59dee

Arm Limited and Contributors

May 31, 2021

CONTENTS

Table 1: Revision History

Date	Issue	Changes
20 Sep 2017	0.51	• Confidentiality Change, EBBR version 0.51
12 Jul 2018	0.6	• Relicense to CC-BY-SA 4.0 • Added Devicetree requirements • Added Multiprocessor boot requirements • Transitioned to reStructuredText and GitHub • Added firmware on shared media requirements • RTC is optional • Add constraints on sharing devices between firmware and OS • Add large note on implementation of runtime modification of non-volatile variables
18 Oct 2018	0.7	• Add AArch32 details • Refactor Runtime Services text after face to fact meeting at Linaro Connect YVR18
12 Mar 2019	0.8	• Update language around SetVariable() and what is available during runtime services • Editorial changes preparing for v1.0
31 Mar 2019	1.0	• Remove unnecessary UEFI requirements appendix • Allow for ACPI vendor id in firmware path
5 Aug 2020	1.0.1	• Update to UEFI 2.8 Errata A • Specify UUID for passing DTB • Typo and editorial fixes • Document the release process
23 Apr 2021	2.0.0	• Reduce the number of UEFI required elements needed for compliance. • Add requirement for UpdateCapsule() runtime service. • Updated firmware shared storage requirements • Refined RTC requirements • Fixed ResetSystem() to correctly describe failure condition

ABOUT THIS DOCUMENT

1.1 Introduction

This Embedded Base Boot Requirements (EBBR) specification defines an interface between platform firmware and an operating system that is suitable for embedded platforms. EBBR compliant platforms present a consistent interface that will boot an EBBR compliant operating system without any custom tailoring required. For example, an Arm A-class embedded platform will benefit from a standard interface that supports features such as secure boot and firmware update.

This specification defines the base firmware requirements for EBBR compliant platforms. The requirements in this specification are expected to be minimal yet complete, while leaving plenty of room for innovations and design details. This specification is intended to be OS-neutral.

It leverages the prevalent industry standard firmware specification of [UEFI].

Comments or change requests can be sent to *boot-architecture@lists.linaro.org*.

1.2 Guiding Principles

EBBR as a specification defines requirements on platforms and operating systems, but requirements alone don't provide insight into why the specification is written the way it is, or what problems it is intended to solve. Using the assumption that better understanding of the thought process behind EBBR will result in better implementations, this section is a discussion of the goals and guiding principle that shaped EBBR.

This section should be considered commentary, and not a formal part of the specification.

EBBR was written as a response to the lack of boot sequence standardization in the embedded system ecosystem. As embedded systems are becoming more sophisticated and connected, it is becoming increasingly important for embedded systems to run standard OS distributions and software stacks, or to have consistent behaviour across a large deployment of heterogeneous platforms. However, the lack of consistency between platforms often requires per-platform customization to get an OS image to boot on multiple platforms.

A large part of this ecosystem is based on U-Boot and Linux. Vendors have heavy investments in both projects and are not interested in large scale changes to their firmware architecture. The challenge for EBBR is to define a set of boot standards that reduce the amount of custom engineering required, make it possible for OS distributions to support embedded platforms, while still preserving the firmware stack that product vendors are comfortable with. Or in simpler terms, EBBR is designed to solve the embedded boot mess by adding a defined standard (UEFI) to the existing firmware projects (U-Boot).

However, EBBR is a specification, not an implementation. The goal of EBBR is not to mandate U-Boot and Linux. Rather, it is to mandate interfaces that can be implemented by any firmware or OS project, while at the same time work with both Tianocore/EDK2 and U-Boot to ensure that the EBBR requirements are implemented by both projects.[1]

The following guiding principles are used while developing the EBBR specification.

---

[1] Tianocore/EDK2 and U-Boot are highlighted here because at the time of writing these are the two most important firmware projects that implement UEFI. Tianocore/EDK2 is a full featured UEFI implementation and so should automatically be EBBR compliant. U-Boot is the incumbent firmware project for embedded platforms and has steadily been adding UEFI compliance since 2016.

- Be agnostic about ACPI and Devicetree.

  EBBR explicitly does not require a specific system description language. Both Devicetree and ACPI are supported. The Linux kernel supports both equally well, and so EBBR doesn't require one over the other. However, EBBR does require the system description to be supplied by the platform, not the OS. The platform must also conform to the relevant ACPI or DT specifications and adhere to platform compatibility rules.[2]

- Focus on the UEFI interface, not a specific codebase

  EBBR does not require a specific firmware implementation. Any firmware project can implement these interfaces. Neither U-Boot nor Tianocore/EDK2 are required.

- Design to be implementable and useful today

  The drafting process for EBBR worked closely with U-Boot and Tianocore developers to ensure that current upstream code will meet the requirements.

- Design to be OS independent

  This document uses Linux as an example but other OS's support EBBR compliant systems as well (e.g. FreeBSD, OpenBSD).

- Support multiple architectures

  Any architecture can implement the EBBR requirements. Architecture specific requirements will clearly marked as to which architecture(s) they apply.

- Design for common embedded hardware

  EBBR support will be implemented on existing developer hardware. Generally anything that has a near-upstream U-Boot implementation should be able to implement the EBBR requirements. EBBR was drafted with readily available hardware in mind, like the Raspberry Pi and BeagleBone families of boards, and it is applicable for low cost boards (<$10).

- Plan to evolve over time

  The current release of EBBR is firmly targeted at existing platforms so that gaining EBBR compliance may require a firmware update, but will not require hardware changes for the majority of platforms.

  Future EBBR releases will tighten requirements to add features and improve compatibility, which may affect hardware design choices. However, EBBR will not retroactively revoke support from previously compliant platforms. Instead, new requirements will be clearly documented as being over and above what was required by a previous release. Existing platforms will be able to retain compliance with a previous requirement level. In turn, OS projects and end users can choose what level of EBBR compliance is required for their use case.

1.3 Scope

This document defines a subset of the boot and runtime services, protocols and configuration tables defined in the UEFI specification [UEFI] that is provided to an Operating System or hypervisor.

This specification defines the boot and runtime services for a physical system, including services that are required for virtualization. It does not define a standardized abstract virtual machine view for a Guest Operating System.

This specification is referenced by the Arm Base Boot Requirements Specification [ArmBBR] § 4.3. The UEFI requirements found in this document are similar but not identical to the requirements found in BBR. EBBR provides greater flexibility for support embedded designs which cannot easily meet the stricter BBR requirements.

By definition, all BBR compliant systems are also EBBR compliant, but the converse is not true.

---

[2] It must be acknowledged that at the time of writing this document, platform compatibility rules for DT platforms are not well defined or documented. We the authors recognize that this is a problem and are working to solve it in parallel with this specification.

1.4 Conventions Used in this Document

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in **RFC 2119**.

1.5 Cross References

This document cross-references sources that are listed in the References section by using the section sign §.

Examples:

UEFI § 6.1 - Reference to the UEFI specification [UEFI] section 6.1

1.6 Terms and abbreviations

This document uses the following terms and abbreviations.

A64 The 64-bit Arm instruction set used in AArch64 state. All A64 instructions are 32 bits.

AArch32 Arm 32-bit architectures. AArch32 is a roll up term referring to all 32-bit versions of the Arm architecture starting at ARMv4.

AArch64 state The Arm 64-bit Execution state that uses 64-bit general purpose registers, and a 64-bit program counter (PC), Stack Pointer (SP), and exception link registers (ELR).

AArch64 Execution state provides a single instruction set, A64.

EFI Loaded Image An executable image to be run under the UEFI environment, and which uses boot time services.

EL0 The lowest Exception level on AArch64. The Exception level that is used to execute user applications, in Non-secure state.

EL1 Privileged Exception level on AArch64. The Exception level that is used to execute Operating Systems, in Non-secure state.

EL2 Hypervisor Exception level on AArch64. The Exception level that is used to execute hypervisor code. EL2 is always in Non-secure state.

EL3 Secure Monitor Exception level on AArch64. The Exception level that is used to execute Secure Monitor code, which handles the transitions between Non-secure and Secure states. EL3 is always in Secure state.

Logical Unit (LU) A logical unit (LU) is an externally addressable, independent entity within a device. In the context of storage, a single device may use logical units to provide multiple independent storage areas.

OEM Original Equipment Manufacturer. In this document, the final device manufacturer.

SiP Silicon Partner. In this document, the silicon manufacturer.

UEFI Unified Extensible Firmware Interface.

UEFI Boot Services Functionality that is provided to UEFI Loaded Images during the UEFI boot process.

UEFI Runtime Services Functionality that is provided to an Operating System after the ExitBootServices() call.

UEFI

This chapter discusses specific UEFI implementation details for EBBR compliant platforms.

2.1 UEFI Version

This document uses version 2.8 Errata A of the UEFI specification [UEFI].

2.2 UEFI Compliance

EBBR compliant platform shall conform to a subset of the [UEFI] spec as listed in this section. Normally, UEFI compliance would require full compliance with all items listed in UEFI § 2.6. However, the EBBR target market has a reduced set of requirements, and so some UEFI features are omitted as unnecessary.

2.2.1 Required Elements

This section replaces the list of required elements in [UEFI] § 2.6.1. All of the following UEFI elements are required for EBBR compliance.

Table 2.1: UEFI Required Elements

Element	Requirement
EFI_SYSTEM_TABLE	The system table is required to provide required to access UEFI Boot Services, UEFI Runtime Services, consoles, and other firmware, vendor and platform information.
EFI_BOOT_SERVICES	All functions defined as boot services must exist. Methods for unsupported or unimplemented behaviour must return an appropriate error code.
EFI_RUNTIME_SERVICES	All functions defined as runtime services must exist. Methods for unsupported or unimplemented behaviour must return an appropriate error code. If any runtime service is unimplemented, it must be indicated via the *EFI_RT_PROPERTIES_TABLE*.
EFI_LOADED_IMAGE_PROTOCOL	Must be installed for each loaded image
EFI_LOADED_IMAGE_DEVICE_PATH_PROTOCOL	Must be installed for each loaded image
EFI_DEVICE_PATH_PROTOCOL	An *EFI_DEVICE_PATH_PROTOCOL* must be installed onto all device handles provided by the firmware.
EFI_DEVICE_PATH_UTILITIES_PROTOCOL	Interface for creating and manipulating UEFI device paths

Table 2.2: Notable omissions from UEFI § 2.6.1

Element	Note
EFI_DECOMPRESS_PROTOCOL	Native EFI decompression is rarely used and therefore not required.

2.2.2 Required Platform Specific Elements

This section replaces the list of required elements in [UEFI] § 2.6.2. All of the following UEFI elements are required for EBBR compliance.

Table 2.3: UEFI Platform-Specific Required Elements

Element	Description
Console devices	The platform must have at least one console device
EFI_SIMPLE_TEXT_INPUT_PROTOCOL	Needed for console input
EFI_SIMPLE_TEXT_INPUT_EX_PROTOCOL	Needed for console input
EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL	Needed for console output
EFI_DEVICE_PATH_TO_TEXT_PROTOCOL	Needed for console output
EFI_HII_STRING_PROTOCOL	Required by EFI shell and for compliance testing
EFI_HII_DATABASE_PROTOCOL	Required by EFI shell and for compliance testing
EFI_UNICODE_COLLATION2_PROTOCOL	Required by EFI shell and for compliance testing
EFI_BLOCK_IO_PROTOCOL	Required for block device access
EFI_SIMPLE_FILE_SYSTEM_PROTOCOL	Required if booting from block device is supported
EFI_RNG_PROTOCOL	Required if the platform has a hardware entropy source
EFI_SIMPLE_NETWORK_PROTOCOL	Required if the platform has a network device.
HTTP Boot (UEFI § 24.7)	Required if the platform supports network booting

The following table is a list of notable deviations from UEFI § 2.6.2. Many of these deviations are because the EBBR use cases do not require interface specific UEFI protocols, and so they have been made optional.

Table 2.4: Notable Deviations from UEFI § 2.6.2

Element	Description of deviation
LoadImage()	The *LoadImage()* boot service is not required to install an *EFI_HII_PACKAGE_LIST_PROTOCOL* for an image containing a custom PE/COFF resource with the type 'HII'. HII resource images are not needed to run the UEFI shell or the SCT.
ConnectController()	The *ConnectController()* boot service is not required to support the *EFI_PLATFORM_DRIVER_OVERRIDE_PROTOCOL*, *EFI_DRIVER_FAMILY_OVERRIDE_PROTOCOL*, and *EFI_BUS_SPECIFIC_DRIVER_OVERRIDE_PROTOCOL*. These override protocols are only useful if drivers are loaded as EFI binaries by the firmware.
EFI_HII_CONFIG_ACCESS_PROTOCOL	UEFI requires this for console devices, but it is rarely necessary in practice. Therefore this protocol is not required.
EFI_HII_CONFIG_ROUTING_PROTOCOL	UEFI requires this for console devices, but it is rarely necessary in practice. Therefore this protocol is not required.
Graphical console	Platforms with a graphical device are not required to expose it as a graphical console.
EFI_DISK_IO_PROTOCOL	Rarely used interface that isn't required for EBBR use cases
EFI_PXE_BASE_CODE_PROTOCOL	Booting via the Preboot Execution Environment (PXE) is insecure. Loading via PXE is typically executed before launching the first UEFI application.
Network protocols	A full implementation of the UEFI general purpose networking ABIs is not required, including *EFI_NETWORK_INTERFACE_IDENTIFIER_PROTOCOL*, *EFI_MANAGED_NETWORK_PROTOCOL*, *EFI_*_SERVICE_BINDING_PROTOCOL*, or any of the IPv4 or IPv6 protocols.
Byte stream device support (UART)	UEFI protocols not required
PCI bus support	UEFI protocols not required
USB bus support	UEFI protocols not required
NVMe pass through support	UEFI protocols not required
SCSI pass through support	UEFI protocols not required
EFI_DRIVER_FAMILY_OVERRIDE_PROTOCOL	Not required
Option ROM support	In many EBBR use cases there is no requirement to generically support any PCIe add in card at the firmware level. When PCIe devices are used, drivers for the device are often built into the firmware itself rather than loaded as option ROMs. For this reason EBBR implementations are not required to support option ROM loading.

2.2.3 Required Global Variables

EBBR compliant platforms are required to support the following Global Variables as found in [UEFI] § 3.3.

Table 2.5: Required UEFI Variables

Variable Name	Description
Boot####	A boot load option. #### is a numerical hex value
BootCurrent	The boot option that was selected for the current boot
BootNext	The boot option that will be used for the next boot only
BootOrder	An ordered list of boot options. Firmware will try BootNext and each Boot#### entry in the order given by BootOrder to find the first bootable image.
OsIndications	Method for OS to request features from firmware
OsIndicationsSupported	Variable for firmware to indicate which features can be enabled

2.2.4 Block device partitioning

The system firmware must implement support for MBR, GPT and El Torito partitioning on block devices. System firmware may also implement other partitioning methods as needed by the platform, but OS support for other methods is outside the scope of this specification.

2.3 UEFI System Environment and Configuration

The resident UEFI boot-time environment shall use the highest non-secure privilege level available. The exact meaning of this is architecture dependent, as detailed below.

Resident UEFI firmware might target a specific privilege level. In contrast, UEFI Loaded Images, such as third-party drivers and boot applications, must not contain any built-in assumptions that they are to be loaded at a given privilege level during boot time since they can, for example, legitimately be loaded into either EL1 or EL2 on AArch64.

2.3.1 AArch64 Exception Levels

On AArch64 UEFI shall execute as 64-bit code at either EL1 or EL2, depending on whether or not virtualization is available at OS load time.

UEFI Boot at EL2

Most systems are expected to boot UEFI at EL2, to allow for the installation of a hypervisor or a virtualization aware Operating System.

UEFI Boot at EL1

Booting of UEFI at EL1 is most likely employed within a hypervisor hosted Guest Operating System environment, to allow the subsequent booting of a UEFI-compliant Operating System. In this instance, the UEFI boot-time environment can be provided, as a virtualized service, by the hypervisor and not as part of the host firmware.

2.4 UEFI Boot Services

2.4.1 Memory Map

The UEFI environment must provide a system memory map, which must include all appropriate devices and memories that are required for booting and system configuration.

All RAM defined by the UEFI memory map must be identity-mapped, which means that virtual addresses must equal physical addresses.

The default RAM allocated attribute must be *EFI_MEMORY_WB*.

2.4.2 Configuration Tables

A UEFI system that complies with this specification may provide additional tables via the EFI Configuration Table.

Compliant systems are required to provide one, but not both, of the following tables:

- an Advanced Configuration and Power Interface [ACPI] table, or

- a Devicetree [DTSPEC] system description

EBBR systems must not provide both ACPI and Devicetree tables at the same time. Systems that support both interfaces must provide a configuration mechanism to select either ACPI or Devicetree, and must ensure only the selected interface is provided to the OS loader.

Devicetree

If firmware provides a Devicetree system description then it must be provided in Flattened Devicetree Blob (DTB) format version 17 or higher as described in [DTSPEC] § 5.1. The following GUID must be used in the EFI system table ([UEFI] § 4) to identify the DTB. The DTB must be contained in memory of type EfiACPIReclaimMemory. EfiACPIReclaimMemory was chosen to match the recommendation for ACPI tables which fulfill the same task as the DTB.

```
#define EFI_DTB_GUID \
    EFI_GUID(0xb1b621d5, 0xf19c, 0x41a5, \
            0x83, 0x0b, 0xd9, 0x15, 0x2c, 0x69, 0xaa, 0xe0)
```

Firmware must have the DTB resident in memory and installed in the EFI system table before executing any UEFI applications or drivers that are not part of the system firmware image. Once the DTB is installed as a configuration table, the system firmware must not make any modification to it or reference any data contained within the DTB.

UEFI applications are permitted to modify or replace the loaded DTB. System firmware must not depend on any data contained within the DTB. If system firmware makes use of a DTB for its own configuration, it should use a separate private copy that is not installed in the EFI System Table or otherwise be exposed to EFI applications.

2.4.3 UEFI Secure Boot (Optional)

UEFI Secure Boot is optional for this specification.

If Secure Boot is implemented, it must conform to the UEFI specification for Secure Boot. There are no additional requirements for Secure Boot.

2.5 UEFI Runtime Services

UEFI runtime services exist after the call to *ExitBootServices()* and are designed to provide a limited set of persistent services to the platform Operating System or hypervisor. Functions contained in *EFI_RUNTIME_SERVICES* are expected to be available during both boot services and runtime services. However, it isn't always practical for all *EFI_RUNTIME_SERVICES* functions to be callable during runtime services due to hardware limitations. If any *EFI_RUNTIME_SERVICES* functions are only available during boot services then firmware shall provide the *EFI_RT_PROPERTIES_TABLE* to indicate which functions are available during runtime services. Functions that are not available during runtime services shall return *EFI_UNSUPPORTED*.

Table 2.6 details which *EFI_RUNTIME_SERVICES* are required to be implemented during boot services and runtime services.

Table 2.6: *EFI_RUNTIME_SERVICES* Implementation Requirements

EFI_RUNTIME_SERVICES function	Before *ExitBootServices()*	After *ExitBootServices()*
GetTime	Required if RTC present	Optional
SetTime	Required if RTC present	Optional
GetWakeupTime	Required if wakeup supported	Optional
SetWakeupTime	Required if wakeup supported	Optional
SetVirtualAddressMap	N/A	Required
ConvertPointer	N/A	Required
GetVariable	Required	Optional
GetNextVeriableName	Required	Optional
SetVariable	Required	Optional
GetNextHighMonotonicCount	N/A	Optional
ResetSystem	Required	Optional
UpdateCapsule	Required for in-band update	Optional
QueryCapsuleCapabilities	Optional	Optional
QueryVariableInfo	Optional	Optional

2.5.1 Runtime Device Mappings

Firmware shall not create runtime mappings, or perform any runtime IO that will conflict with device access by the OS. Normally this means a device may be controlled by firmware, or controlled by the OS, but not both. E.g. if firmware attempts to access an eMMC device at runtime then it will conflict with transactions being performed by the OS.

Devices that are provided to the OS (i.e., via PCIe discovery or ACPI/DT description) shall not be accessed by firmware at runtime. Similarly, devices retained by firmware (i.e., not discoverable by the OS) shall not be accessed by the OS.

Only devices that explicitly support concurrent access by both firmware and an OS may be mapped at runtime by both firmware and the OS.

Real-time Clock (RTC)

Not all embedded systems include an RTC, and even if one is present, it may not be possible to access the RTC from runtime services. e.g., The RTC may be on a shared I2C bus which runtime services cannot access because it will conflict with the OS.

If an RTC is present, then *GetTime()* and *SetTime()* must be supported before *ExitBootServices()* is called.

However, if firmware does not support access to the RTC after *ExitBootServices()*, then *GetTime()* and *SetTime()* shall return *EFI_UNSUPPORTED* and the OS must use a device driver to control the RTC.

2.5.2 UEFI Reset and Shutdown

ResetSystem() is required to be implemented in boot services, but it is optional for runtime services. During runtime services, the operating system should first attempt to use *ResetSystem()* to reset the system.

If firmware doesn't support *ResetSystem()* during runtime services, then the call will immediately return, and the OS should fall back to an architecture or platform specific reset mechanism.

On AArch64 platforms implementing [PSCI], if *ResetSystem()* is not implemented then the Operating System should fall back to making a PSCI call to reset or shutdown the system.

2.5.3 Runtime Variable Access

There are many platforms where it is difficult to implement *SetVariable()* for non-volatile variables during runtime services because the firmware cannot access storage after *ExitBootServices()* is called.

e.g., If firmware accesses an eMMC device directly at runtime, it will collide with transactions initiated by the OS. Neither U-Boot nor Tianocore have a generic solution for accessing or updating variables stored on shared media.[1]

If a platform does not implement modifying non-volatile variables with *SetVariable()* after *ExitBootServices()*, then firmware shall return *EFI_UNSUPPORTED* for any call to *SetVariable()*, and must advertise that *SetVariable()* isn't available during runtime services via the *RuntimeServicesSupported* value in the *EFI_RT_PROPERTIES_TABLE* as defined in [UEFI] § 4.6. EFI applications can read *RuntimeServicesSupported* to determine if calls to *SetVariable()* need to be performed before calling *ExitBootServices()*.

Even when *SetVariable()* is not supported during runtime services, firmware should cache variable names and values in EfiRuntimeServicesData memory so that *GetVariable()* and *GetNextVeriableName()* can behave as specified.

2.5.4 Firmware Update

Being able to update firmware to address security issues is a key feature of secure platforms. EBBR platforms are required to implement either an in-band or an out-of-band firmware update mechanism.

If firmware update is performed in-band (firmware on the application processor updates itself), then the firmware shall implement the *UpdateCapsule()* runtime service and accept updates in the "Firmware Management Protocol Data Capsule Structure" format as described in [UEFI] § 23.3, "Delivering Capsules Containing Updates to Firmware Management Protocol.[2] Firmware is also required to provide an EFI System Resource Table (ESRT). [UEFI] § 23.4 Every firmware image that can be updated in-band must be described in the ESRT.

If firmware update is performed out-of-band (e.g., by an independent Baseboard Management Controller (BMC), or firmware is provided by a hypervisor), then the platform is not required to implement the *UpdateCapsule()* runtime service.

UpdateCapsule() is only required before *ExitBootServices()* is called.

---

[1] It is worth noting that OP-TEE has a similar problem regarding secure storage. OP-TEE's chosen solution is to rely on an OS supplicant agent to perform storage operations on behalf of OP-TEE. The same solution may be applicable to solving the UEFI non-volatile variable problem, but it requires additional OS support to work. Regardless, EBBR compliance does not require *SetVariable()* support during runtime services.

https://optee.readthedocs.io/en/latest/architecture/secure_storage.html

[2] The *UpdateCapsule()* runtime service is expected to be suitable for use by generic firmware update services like fwupd and Windows Update. Both fwupd and Windows Update read the ESRT table to determine what firmware can be updated, and use an EFI helper application to call *UpdateCapsule()* before *ExitBootServices()* is called.

https://fwupd.org/

PRIVILEGED OR SECURE FIRMWARE

3.1 AArch32 Multiprocessor Startup Protocol

There is no standard multiprocessor startup or CPU power management mechanism for ARMv7 and earlier platforms. The OS is expected to use platform specific drivers for CPU power management. Firmware must advertize the CPU power management mechanism in the Devicetree system description or the ACPI tables so that the OS can enable the correct driver. At ExitBootServices() time, all secondary CPUs must be parked or powered off.

3.2 AArch64 Multiprocessor Startup Protocol

On AArch64 platforms, Firmware resident in Trustzone EL3 must implement and conform to the Power State Coordination Interface specification [PSCI].

Platforms without EL3 must implement one of:

- PSCI at EL2 (leaving only EL1 available to an operating system)
- Linux AArch64 spin tables [LINUXA64BOOT] (Devicetree only)

However, the spin table protocol is strongly discouraged. Future versions of this specification will only allow PSCI, and PSCI should be implemented in all new designs.

FIRMWARE STORAGE

In general, EBBR compliant platforms should use dedicated storage for boot firmware images and data, independent of the storage used for OS partitions and the EFI System Partition (ESP). This could be a physically separate device (e.g. SPI flash), or a dedicated logical unit (LU) within a device (e.g. eMMC boot partition,[1] or UFS boot LU[2]).

However, many embedded systems have size, cost, or implementation constraints that make separate firmware storage unfeasible. On such systems, firmware and the OS reside in the same storage device. Care must be taken to ensure firmware kept in normal storage does not conflict with normal usage of the media by an OS.

- Firmware must be stored on the media in a way that does not conflict with normal partitioning and usage by the operating system.

- Normal operation of the OS must not interfere with firmware files.

- Firmware needs a method to modify variable storage at runtime while the OS controls access to the device.[3]

4.1 Partitioning of Shared Storage

The shared storage device must use the GUID Partition Table (GPT) disk layout as defined in [UEFI] § 5.3, unless the platform boot sequence is fundamentally incompatible with the GPT disk layout. In which case, a legacy Master Boot Recored (MBR) must be used.[4]

Warning: MBR partitioning is deprecated and only included for legacy support. All new platforms are expected to use GPT partitioning. GPT partitioning supports a much larger number of partitions, and has built in resiliency.

A future version of this specification will disallow the use of MBR partitioning.

Firmware images and data in shared storage should be contained in partitions described by the GPT or MBR. The platform should locate firmware by searching the partition table for the partition(s) containing firmware.

However, some SoCs load firmware from a fixed offset into the storage media. In this case, to protect against partitioning tools overwriting firmware, the partition table must be formed in a way to protect the firmware image(s) as described in sections *GPT partitioning* and *MBR partitioning*.

---

[1] Watch out for the ambiguity of the word 'partition'. In most of this document, a 'partition' is a contiguous region of a block device as described by a GPT or MBR partition table, but eMMC devices also provide a dedicated 'boot partition' that is addressed separately from the main storage region, and does not appear in the partition table.

[2] For the purposes of this document, logical units are treated as independent storage devices, each with their own GPT or MBR partition table. A platform that uses one LU for firmware, and another LU for OS partitions and the ESP is considered to be using dedicated firmware storage.

[3] Runtime access to firmware data may still be an issue when firmware is stored in a dedicated LU, simply because the OS remains in control of the storage device command stream. If firmware doesn't have a dedicated channel to the storage device, then the OS must proxy all runtime storage IO.

[4] For example, if the SoC boot ROM requires an MBR to find the next stage firmware image, then it is incompatible with the GPT boot layout. Similarly, if the boot ROM expects the next stage firmware to be located at LBA1 (the location of the GPT Header), then it is incompatible with the GPT disk layout. In both cases the shared storage device must use legacy MBR partitioning.

Automatic partitioning tools (e.g. an OS installer) must not delete the protective information in the partition table, or delete, move, or modify protective partition entries. Manual partitioning tools should provide warnings when modifying protective partitions.

Warning: Fixed offsets to firmware data is supported only for legacy reasons. All new platforms are expected to use partitions to locate firmware files.

A future version of this specification will disallow the use of fixed offsets.

4.1.1 GPT partitioning

The partition table must strictly conform to the UEFI specification and include a protective MBR authored exactly as described in [UEFI] § 5.3 (hybrid partitioning schemes are not permitted).

Fixed-location firmware images must be protected by creating protective partition entries, or by placing GPT data structures away from the LBAs occupied by firmware,

Protective partitions are entries in the partition table that cover the LBA region occupied by firmware and have the 'Required Partition' attribute set. A protective partition must use a *PartitionTypeGUID* that identifies it as a firmware protective partition. (e.g., don't reuse a GUID used by non-protective partitions). There are no requirements on the contents or layout of the firmware protective partition.

Placing GPT data structures away from firmware images can be accomplished by adjusting the GUID Partition Entry array location (adjusting the values of *PartitionEntryLBA* and *NumberOfPartitionEntries*, and *SizeOfPartitionEntry*), or by specifying the usable LBAs (Choosing *FirstUsableLBA/LastUsableLBA* to not overlap the fixed firmware location). See [UEFI] § 5.3.2.

Given the choice, platforms should use protective partitions over adjusting the placement of GPT data structures because protective partitions provide explicit information about the protected region.

MBR partitioning

If firmware is at a fixed location entirely within the first 1MiB of storage (<= LBA2047) then no protective partitions are required. If firmware resides in a fixed location outside the first 1MiB, then a protective partition must be used to cover the firmware LBAs. Protective partitions should have a partition type of 0xF8 unless an immutable feature of the platform makes this impossible.

OS partitioning tools must not create partitions in the first 1MiB of the storage device, and must not remove protective partitions.

4.2 Firmware Partition Filesystem

Where possible, firmware images and data should be stored in a filesystem. Firmware can be stored either in a dedicated firmware partition, or in certain circumstances in the UEFI System Partition (ESP). Using a filesystem makes it simpler to manage multiple firmware files and makes it possible for a single disk image to contain firmware for multiple platforms.

When firmware is stored in the ESP, the ESP should contain a directory named /FIRMWARE in the root directory, and all firmware images and data should be stored in platform vendor subdirectories under /FIRMWARE.

Dedicated firmware partitions should be formatted with a FAT filesystem as defined by the UEFI specification. Dedicated firmware partitions should use the same /FIRMWARE directory hierarchy. OS tools shall ignore dedicated firmware partitions, and shall not attempt to use a dedicated firmware partition as an ESP.

Vendors may choose their own subdirectory name under /FIRMWARE, but shall choose names that do not conflict with other vendors. Normally the vendor name will be the name of the SoC vendor, because the firmware directory name will be hard coded in the SoC's boot ROM. Vendors are recommended to use their Devicetree vendor prefix or ACPI vendor ID as their vendor subdirectory name.

Vendors are free to decide how to structure subdirectories under their own vendor directory, but they shall use a naming convention that allows multiple SoCs to be supported in the same filesystem.

For example, a vendor named Acme with two SoCs, AM100 & AM300, could choose to use the SoC part number as a subdirectory in the firmware path:

```
/FIRMWARE
  /ACME
    /AM100
      fw.img
    /AM300
      fw.img
```

It is also recommended for dedicated firmware partitions to use the /FIRMWARE file hierarchy.

The following is a sample directory structure for firmware files:

```
/FIRMWARE
  /<Vendor 1 Directory>
    /<SoC A Directory>
        <Firmware image>
        <Firmware data>
    /<SoC B Directory>
        <Firmware image>
        <Firmware data>
  /<Vendor 2 Directory>
    <Common Firmware image>
    <Common Firmware data>
  /<Vendor 3 Directory>
    /<SoC E Directory>
        <Firmware image>
```

Operating systems and installers should not manipulate any files in the /FIRMWARE hierarchy during normal operation.

The sections below discuss the requirements when using both fixed and removable storage. However, it should be noted that the recommended behaviour of firmware should be identical regardless of storage type. In both cases, the recommended boot sequence is to first search for firmware in a dedicated firmware partition, and second search for firmware in the ESP. The only difference between fixed and removable storage is the recommended factory settings for the platform.

4.2.1 Fixed Shared Storage

Fixed storage is storage that is permanently attached to the platform, and cannot be moved between systems. eMMC and Universal Flash Storage (UFS) device are often used as shared fixed storage for both firmware and the OS.

Where possible, it is preferred for the system to boot from a dedicated boot region on media that provides one (e.g., eMMC) that is sufficiently large. Otherwise, the platform storage should be pre-formatted in the factory with a partition table, a dedicated firmware partition, and firmware binaries installed.

Operating systems must not use the dedicated firmware partition for installing EFI applications including, but not limited to, the OS loader and OS specific files. Instead, a normal ESP should be created. OS partitioning tools must take care not to modify or delete dedicated firmware partitions.

4.2.2 Removable Shared Storage

Removable storage is any media that can be physically removed from the system and moved to another machine as part of normal operation (e.g., SD cards, USB thumb drives, and CDs).

There are two primary scenarios for storing firmware on removable media.

1. Platforms that only have removable media (e.g., The Raspberry Pi has an SD card slot, but no fixed storage).

2. Recovery when on-board firmware has been corrupted. If firmware on fixed media has been corrupted, some platforms support loading firmware from removable media which can then be used to recover the platform.

In both cases, it is desirable to start with a stock OS boot image, copy it to the media (SD or USB), and then add the necessary firmware files to make the platform bootable. Typically, OS boot images won't include a dedicated firmware partition, and it is inconvenient to repartition the media to add one. It is simpler and easier for the user if they are able to copy the required firmware files into the /FIRMWARE directory tree on the ESP using the basic file manager tools provided by all desktop operating systems.

On removable media, firmware should be stored in the ESP under the /FIRMWARE directory structure as described in *Firmware Partition Filesystem*. Platform vendors should support their platform by providing a single .zip file that places all the required firmware files in the correct locations when extracted in the ESP /FIRMWARE directory. For simplicity sake, it is expected the same .zip file will recover the firmware files in a dedicated firmware partition.

[ACPI]   Advanced Configuration and Power Interface specification v6.2A, September 2017, UEFI Forum

[DTSPEC]   Devicetree specification v0.3, Devicetree.org

[LINUXA64BOOT] Linux Documentation/arm64/booting.rst, Linux kernel

[PSCI]   Power State Coordination Interface Issue C (PSCI v1.0) 30 January 2015, Arm Limited

[ArmBBR]   Arm Base Boot Requirements specification Issue F (v1.0) 6 Oct 2020, Arm Limited

[UEFI]   Unified Extensable Firmware Interface Specification v2.8 Errata A, February 2020, UEFI Forum